



SILICON VALLEY
DATA SCIENCE



**SPARK
SUMMIT**
2017

Data Wrangling with PySpark for Data Scientists Who Know Pandas

Dr. Andrew Ray



ANDREW RAY

@collegeisfun

Spark contributor


- SparkSQL
- GraphX
- PySpark

Principal Data Engineer @ SVDS

Previously Data Sci. @ Walmart

PhD Mathematics @ Nebraska





Silicon Valley Data Science is a boutique consulting firm focused on transforming your business through data science and engineering.



OUR SERVICES




DATA
STRATEGY



ARCHITECTURE



AGILE
ENGINEERING



AGILE
DATA SCIENCE




COME SEE US & SAY HELLO!

Wednesday, June 7

- Write Graph Algorithms Like a Boss with Andrew Ray (3:20 pm)





To view SVDS speakers and scheduling,
or to receive a copy of our slides, go to:

www.svds.com/SparkSummit17



- Why PySpark?
- Primer
- Setup & Run
- vs. Pandas
- Best Practices

OUTLINE





PySpark

WHY?



A vertical strip on the left side of the slide features a close-up of a lit sparkler. Bright yellow and orange sparks are exploding from the tip, creating a dynamic, starburst-like pattern against a dark background.

WHY?

Do you:

1. Already know Python & Pandas?
2. Love DataFrames?
3. Want to work with Big Data?

Then PySpark is the answer!



WHAT DO I GET WITH PYSPARK?

Gain

- Work with big data
- Native SQL
- Decent documentation

Lose

- Amazing documentation
- Easy plotting
- Indices



PYSPARK

A (quick) Primer



Apache Spark is a fast and general engine
for large-scale data processing.





PRIMER

Distributed compute

- YARN, Mesos, Standalone cluster

Abstractions

- RDD—distributed collection of objects
- Dataframe—distributed dataset of tabular data.
 - Integrated SQL
 - ML Algorithms





IMPORTANT CONCEPTS

Immutable

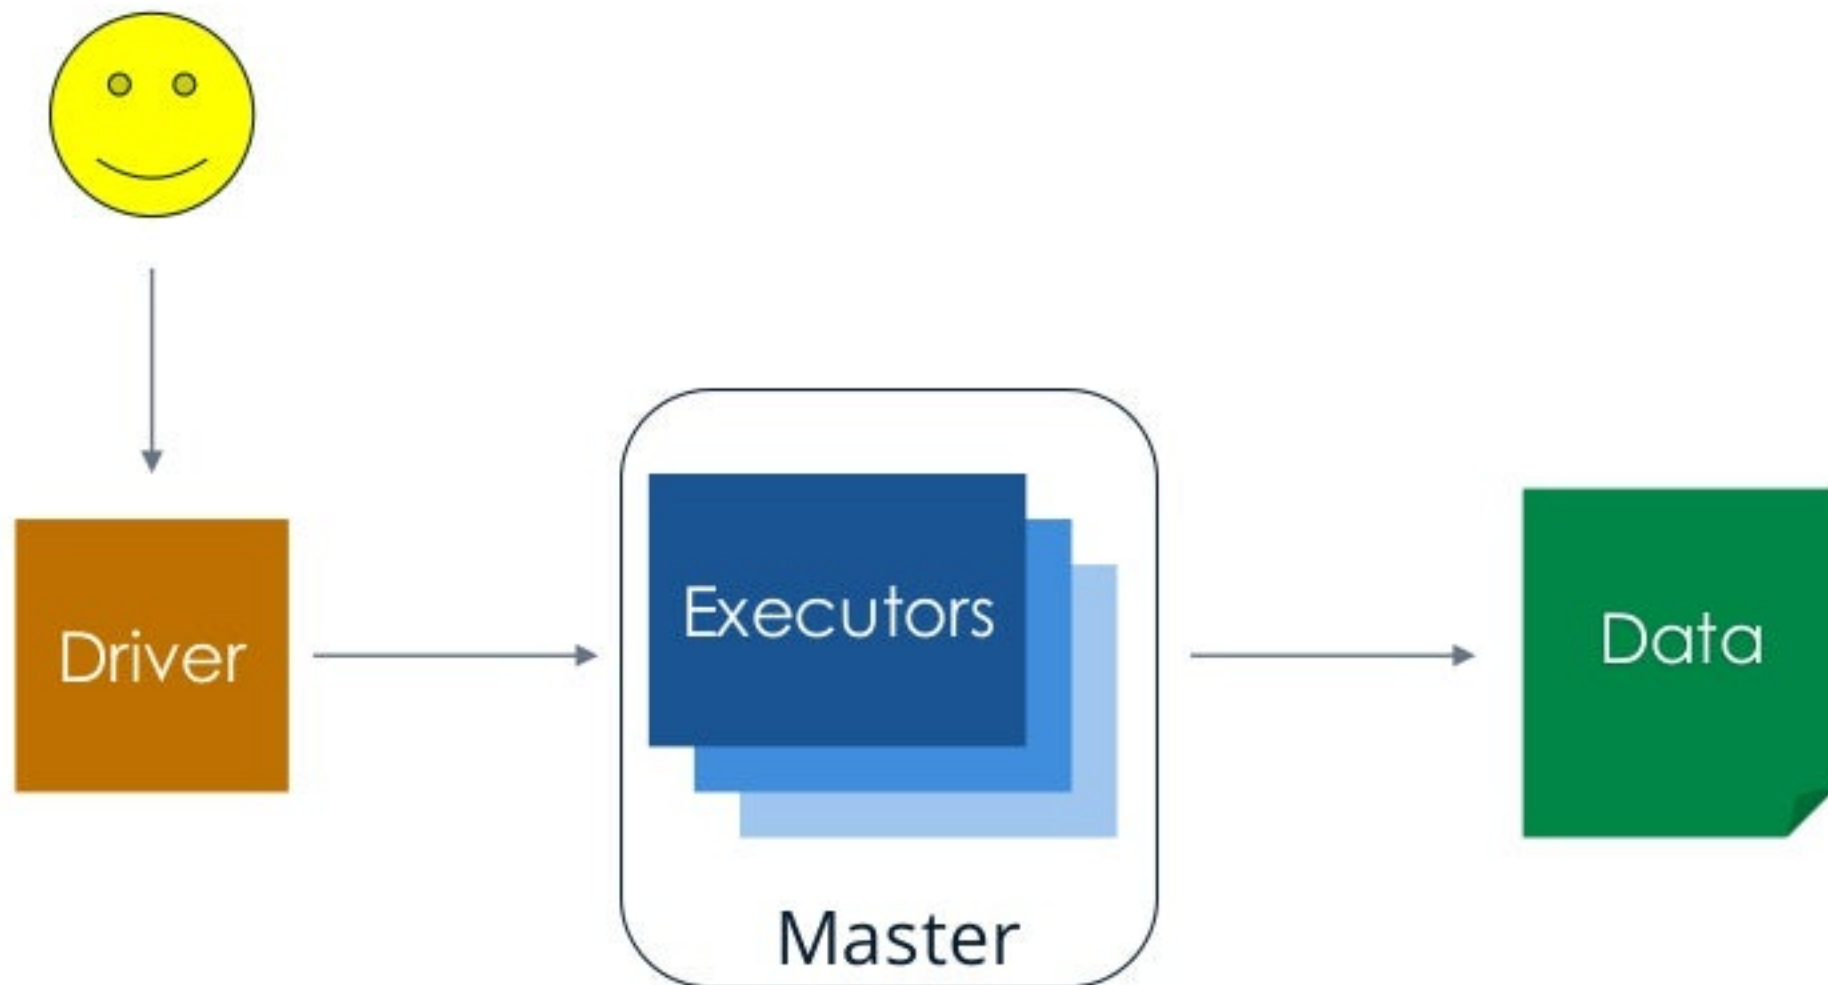
- Changes create new object references
- Old versions are unchanged

Lazy

- Compute does not happen until output is requested



ARCHITECTURE





Setup and Run PYSPARK



SETUP

- Option 1:
 - <https://spark.apache.org/downloads.html>
 - `tar -xzf spark-2.1.1-bin-hadoop2.7.tgz`
 - `PATH="$PATH:$(pwd)/spark-2.1.1-bin-hadoop2.7/bin"`
- Option 2:
 - `conda install -c conda-forge pyspark=2.1.1`
- (Future) Option 3:
 - `pip install pyspark`



RUN

- A. pyspark
- B. PYSPARK_DRIVER_PYTHON=**ipython** pyspark
- C. PYSPARK_DRIVER_PYTHON=**jupyter** \
PYSPARK_DRIVER_PYTHON_OPTS=**notebook** \
pyspark



PYSPARK

vs. Pandas



LOAD CSV

Pandas

```
df = pd.read_csv("mtcars.csv")
```

PySpark

```
df = spark.read.csv("mtcars.csv")
```



LOAD CSV

Pandas

```
df = pd.read_csv("mtcars.csv")
```

PySpark

```
df = spark.read \
    .options(header=True, inferSchema=True) \
    .csv("mtcars.csv")
```



VIEW DATAFRAME

Pandas

```
df  
df.head(10)
```

PySpark

```
df.show()  
df.show(10)
```



COLUMNS AND DATA TYPES

Pandas

```
df.columns  
df.dtypes
```

PySpark

```
df.columns  
df.dtypes
```



RENAME COLUMNS

Pandas

```
df.columns = ['a', 'b', 'c']  
df.rename(columns = {'old': 'new'})
```

PySpark

```
df.toDF('a', 'b', 'c')  
df.withColumnRenamed('old', 'new')
```



DROP COLUMN

Pandas

```
df.drop('mpg', axis=1)
```

PySpark

```
df.drop('mpg')
```



FILTERING

Pandas

```
df[df.mpg < 20]  
df[(df.mpg < 20) & (df.cyl == 6)]
```

PySpark

```
df[df.mpg < 20]  
df[(df.mpg < 20) & (df.cyl == 6)]
```



ADD COLUMN

Pandas

```
df['gpm'] = 1 / df.mpg
```

PySpark

```
df.withColumn('gpm', 1 / df.mpg)
```



FILL NULLS

Pandas

`df.fillna(0)` ← Many more options

PySpark

`df.fillna(0)`



AGGREGATION

Pandas

```
df.groupby(['cyl', 'gear']) \  
    .agg({'mpg': 'mean', 'disp': 'min'})
```

PySpark

```
df.groupby(['cyl', 'gear']) \  
    .agg({'mpg': 'mean', 'disp': 'min'})
```



OK we get the point



STANDARD TRANSFORMATIONS

Pandas

```
import numpy as np  
df['logdisp'] = np.log(df.disp)
```

PySpark

```
import pyspark.sql.functions as F  
df.withColumn('logdisp', F.log(df.disp))
```



KEEP IT IN THE JVM

```
import pyspark.sql.functions as F
```

```
AutoBatchedSerializer collect_set expr length rank substring Column column ctorial  
levenshtein regexp_extract substring_index Dataame concat rst lit regexp_replace sum  
PickleSerializer concat_ws oor locate repeat sumDistinct SparkContext conv rmat_number log  
reverse sys StringType corr rmat_string log10 rint tan UserDenedction cos om_json log1p  
round tanh abs cosh om_unixtime log2 row_number toDegrees acos count om_utc_timestamp  
lower rpad toRadians add_months countDistinct get_json_object lpad rtrim to_date  
approxCountDistinct covar_pop greatest ltrim second to_json approx_count_distinct  
covar_samp grouping map sha1 to_utc_timestamp array crc32 grouping_id math sha2 translate  
array_contains create_map hash max shile trim asc cume_dist hex md5 shiRight trunc ascii  
current_date hour mean shiRightUnsigned udasin current_timestamp hypot min signum unbase64  
atan date_add ignore_unicode_prex minute sin unhex atan2 date_rmat initcap  
monotonically_increasing_id since unix_timestamp avg date_sub input_le_name month sinh  
upper base64 datedi instr months_between size v bin dayoonth isnan nanvl skewness var_pop  
bitwiseNOT dayoear isnull next_day sort_array var_samp blacklist decode json_tuple ntile  
soundex variance broadcast degrees k percent_rank spark_partition_id weekoear bound  
dense_rank kurtosis posexplode split when cbrt desc lag pow sqrt window ceil encode last  
quarter stddev year coalesce exp last_day radians stddev_pop col explode lead rand  
stddev_samp collect_list expm1 least randn struct
```



ROW CONDITIONAL STATEMENTS

Pandas

```
df['cond']=df.apply(lambda r:  
    1 if r.mpg > 20 else 2 if r.cyl == 6 else 3,  
    axis=1)
```

PySpark

```
import pyspark.sql.functions as F  
df.withColumn('cond', \  
    F.when(df.mpg > 20, 1) \  
    .when(df.cyl == 6, 2) \  
    .otherwise(3))
```



PYTHON WHEN REQUIRED

Pandas

```
df['disp1'] = df.disp.apply(lambda x: x+1)
```

PySpark

```
import pyspark.sql.functions as F
from pyspark.sql.types import DoubleType
fn = F.udf(lambda x: x+1, DoubleType())
df.withColumn('disp1', fn(df.disp))
```



MERGE/JOIN DATAFRAMES

Pandas

```
left.merge(right, on='key')  
left.merge(right, left_on='a', right_on='b')
```

PySpark

```
left.join(right, on='key')  
left.join(right, left.a == right.b)
```



PIVOT TABLE

Pandas

```
pd.pivot_table(df, values='D', \
               index=['A', 'B'], columns=['C'], \
               aggfunc=np.sum)
```

PySpark

```
df.groupBy("A", "B").pivot("C").sum("D")
```



SUMMARY STATISTICS

Pandas

```
df.describe()
```

PySpark

```
df.describe().show()    (only count, mean, stddev, min, max)
```

```
df.selectExpr(  
    "percentile_approx(mpg, array(.25, .5, .75)) as mpg"  
).show()
```



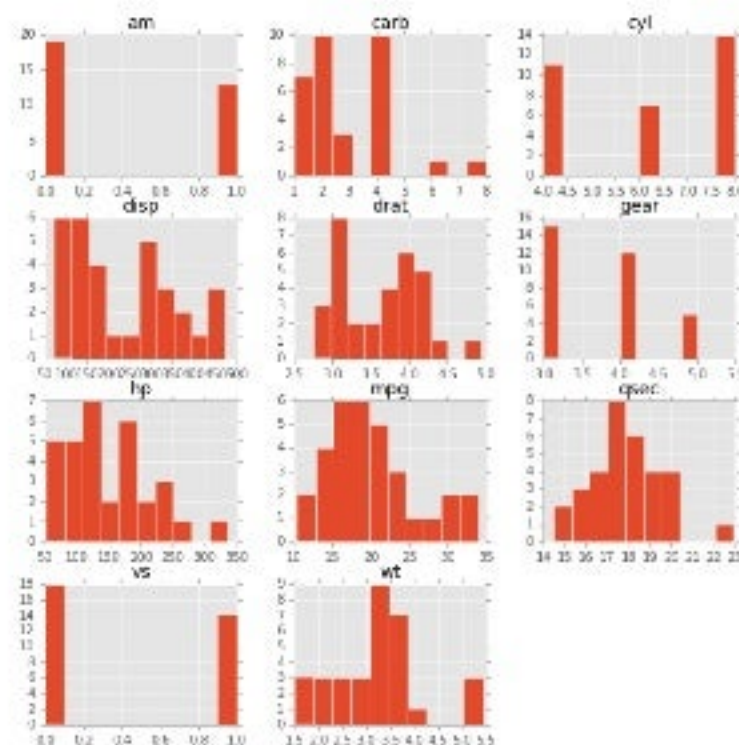
HISTOGRAM

Pandas

```
df.hist()
```

PySpark

```
df.sample(False, 0.1).toPandas().hist()
```



SQL

Pandas

n/a

PySpark

```
df.createOrReplaceTempView('foo')  
df2 = spark.sql('select * from foo')
```



PYSPARK

Best Practices



MAKE SURE TO

- Use `pyspark.sql.functions` and other built in functions.
- Use the same version of python and packages on cluster as driver.
- Check out the UI at <http://localhost:4040/>
- Learn about SSH port forwarding
- Check out Spark MLlib
- RTFM: <https://spark.apache.org/docs/latest/>



THINGS NOT TO DO

- Try to iterate through rows
- Hard code a master in your driver
 - Use `spark-submit` for that
- `df.toPandas().head()`
 - instead do: `df.limit(5).toPandas()`



IF THINGS GO WRONG

- Don't panic!
- Read the error
- Google it
- Search/Ask Stack Overflow (tag [apache-spark](#))
- Search/Ask the user list: [user@spark.apache.org](#)
- Find a bug? Make a JIRA ticket:
<https://issues.apache.org/jira/browse/SPARK/>





SILICON VALLEY
DATA SCIENCE

info@svds.com
@SVDataScience

THANK YOU

Andrew Ray
andrew@svds.com
@collegeisfun